

Maintenance-Constrained Locomotive Assignment: NP-Hardness, Exact Algorithms, and Practical Approximations

Ahmed Rageeb Ahsan, Saiful Islam

University of Florida

Gainesville, Florida, United States of America

ahmedrageebahsan@ufl.edu, saiful.islam@ufl.edu

Abstract

This paper studies the *Maintenance-Constrained Locomotive Assignment* (MCLA) problem, a real-world railway scheduling problem in which trains with different driving requirements must be assigned to locomotives with limited maintenance windows. We show that MCLA is strongly NP-complete by giving a direct polynomial-time reduction from 3-Partition. We provide a formal correctness proof, discuss the optimization version, and show the equivalence to classical Bin Packing. We then present exact exponential-time algorithms (DP over subsets), greedy approximation methods (FFD/BFD), and a compact integer linear programming formulation useful for real scheduling systems. We outline an experimental evaluation framework and provide complete Python implementations of the algorithms. This work connects a real railway resource allocation problem to well-studied NP-hard combinatorial problems, enabling both theoretical analysis and practical solution design.

1 Introduction

Efficient locomotive assignment is a fundamental scheduling challenge in modern railway systems. Locomotives must be assigned to sequences of train services while satisfying technical constraints such as engine wear, inspection intervals, and mandatory maintenance windows. Railway operators must minimize the size of the locomotive fleet while ensuring every train is serviced and every locomotive receives maintenance before exceeding its usage limit.

These constraints create a problem structurally similar to *Bin Packing*, but with a real-world maintenance interpretation. This paper formalizes this problem as the *Maintenance-Constrained Locomotive Assignment* (MCLA) problem, proves it strongly NP-complete, and provides practical algorithms.

We follow the methodology of standard hardness proofs used in real-world optimization problems [1], as well as the locomotive scheduling literature (e.g., Borndörfer et al., Meng et al., and recent optimization theses).

2 Problem Definition

2.1 Real-World Problem

Each train service i requires a_i kilometers of locomotive operation between maintenance visits. A locomotive can operate at most B kilometers before requiring maintenance. The goal is to assign all trains to a fleet of m locomotives.

Maintenance-Constrained Locomotive Assignment (MCLA):

- Each train i has usage $a_i > 0$.
- A locomotive can accumulate at most B usage before going to maintenance.
- A locomotive may service multiple trains, in any order.

Decision version:

Given a_1, \dots, a_n , capacity B , and m locomotives, does there exist an assignment such that no locomotive exceeds B total usage?

Optimization version:

Minimize number of locomotives required.

2.2 Abstract Formulation

Equivalent to:

Partition $\{a_1, \dots, a_n\}$ into at most m subsets G_j such that

$$\sum_{a \in G_j} a \leq B, \quad \forall j.$$

This is literally *bin packing* with capacity B .

2.3 Illustrative Example (Bin Packing Interpretation)

Consider a locomotive with maintenance bound $B = 500$ km and the following train services:

$$a = (220, 180, 150, 140, 130).$$

The Maintenance-Constrained Locomotive Assignment (MCLA) problem asks for a partition of a into the smallest number of subsets G_1, G_2, \dots such that

$$\sum_{i \in G_j} a_i \leq B \quad \text{for all } j.$$

A feasible assignment is:

$$G_1 = \{220, 180\} \quad (\text{sum} = 400),$$

$$G_2 = \{150, 140\} \quad (\text{sum} = 290),$$

$$G_3 = \{130\} \quad (\text{sum} = 130).$$

Thus three locomotives are required. This is exactly the classical bin packing formulation with:

- items = train usages a_i ,
- bins = locomotives,
- bin capacity = maintenance limit B .

The number of ways to group them is enormous (exponential). This is why the problem becomes NP-hard.

3 NP-Completeness Proof

We reduce from **3-Partition**, one of the strongest NP-complete problems.

3.1 3-Partition Problem

Given $3m$ integers x_1, \dots, x_{3m} and integer B such that

$$\frac{B}{4} < x_i < \frac{B}{2} \quad \text{and} \quad \sum_{i=1}^{3m} x_i = mB,$$

decide whether the integers can be partitioned into m triples each summing to B .

This problem is strongly NP-complete.

3.2 Reduction Construction

Given an instance of 3-Partition, construct MCLA instance:

- $n = 3m$ train segments.
- $a_i = x_i$ (direct copy).
- Maintenance capacity B unchanged.
- Number of locomotives m .

Clearly polynomial time.




Figure 1: Reduction from 3-Partition to MCLA.

3.3 Correctness Proof

THEOREM 3.1. 3-Partition is a YES instance if and only if the constructed MCLA instance is YES.

PROOF. (\Rightarrow) Suppose the 3-Partition instance has m triples each summing to B . Assign each triple to one locomotive. Total usage = $B \leq B$. Feasible.

(\Leftarrow) Suppose the MCLA instance is feasible. Total usage is mB . With m locomotives, each of capacity B , every locomotive must have exactly B usage.

Because each a_i satisfies $a_i > B/4$, no locomotive can contain 4 or more items.

Because $a_i < B/2$, no locomotive can contain 1 item whose load is B .

Thus each locomotive contains exactly 3 items, and their sum must be exactly B . These 3-item groups form the required 3-partition. \square

COROLLARY 3.2. MCLA is strongly NP-complete.

4 Algorithms

4.1 Exact Exponential-Time DP

We use subset DP:

$$dp[\text{mask}] = \min \text{ bins used for items in mask.}$$

Transition either places item into current bin or opens a new bin. Time complexity: $O(n2^n)$.

Algorithm 1 Exact Exponential-Time DP for MCLA / Bin Packing

Require: Item sizes a_1, \dots, a_n , capacity B

Ensure: Minimum number of locomotives

```

1:  $N \leftarrow 2^n$ 
2: Initialize  $DP[\text{mask}][c] \leftarrow \infty$  for all mask,  $0 \leq c \leq B$ 
3:  $DP[0][B] \leftarrow 1$             $\triangleright$  Start with one empty locomotive
4: for mask from 0 to  $N - 1$  do
5:   for remaining capacity  $c$  from 0 to  $B$  do
6:     if  $DP[\text{mask}][c] = \infty$  then
7:       continue
8:     end if
9:     for each item  $i$  not in mask do
10:      newMask  $\leftarrow$  mask  $\cup \{i\}$ 
11:      if  $a_i \leq c$  then           $\triangleright$  Place in current locomotive
12:         $DP[\text{newMask}][c - a_i] \leftarrow \min(DP[\text{newMask}][c - a_i], DP[\text{mask}][c])$ 
13:      end if
14:    Open a new locomotive
15:     $DP[\text{newMask}][B - a_i] \leftarrow \min(DP[\text{newMask}][B - a_i], DP[\text{mask}][c] + 1)$ 
16:  end for
17: end for
18: end for
19: return  $\min_{0 \leq c \leq B} DP[(1 \ll n) - 1][c]$ 
  
```

4.2 Greedy Approximation: First-Fit Decreasing (FFD)

Algorithm 2 First-Fit Decreasing

```

1: Sort items in nonincreasing order.
2: Initialize empty list of bins.
3: for each item do
4:   Place into first bin with remaining capacity.
5:   if none then
6:     Open new bin.
7:   end if
8: end for
  
```

FFD satisfies:

$$FFD(I) \leq \left\lceil \frac{11}{9} OPT(I) \right\rceil + 1.$$

4.3 Integer Linear Programming Formulation

Variables:

$$x_{ij} = \begin{cases} 1 & \text{train } i \text{ in locomotive } j \\ 0 & \text{otherwise} \end{cases}, \quad y_j = 1 \text{ if locomotive } j \text{ used.}$$

Constraints:

$$\begin{aligned} \sum_j x_{ij} &= 1 \quad \forall i \\ \sum_i a_i x_{ij} &\leq B y_j \quad \forall j \end{aligned}$$

Minimize:

$$\sum_j y_j.$$

Works well in practice.

4.4 Local Search Heuristics

- pairwise item swaps,
- k-exchange moves,
- re-optimization of small subsets with DP,
- using FFD output as warm start for ILP.

5 Experimental Result

In our experimental evaluation, we implemented the exact exponential-time dynamic programming (DP) algorithm and the greedy First-Fit Decreasing (FFD) heuristic to solve the Maintenance-Constrained Locomotive Assignment problem. The DP algorithm guarantees an optimal solution but has exponential runtime complexity, limiting its practicality to small instances. In contrast, FFD provides fast, near-optimal solutions suitable for larger problem sizes.

We generate synthetic test instances by sampling train usage values uniformly at random between 1 and half of the maintenance capacity B . This avoids trivial cases where one item fits alone or all fit in one bin. The capacity B is fixed at 100 for most experiments. Instance sizes n vary from 8 to 17 trains to keep the exact dynamic programming (DP) algorithm computationally feasible.

Figure 2 compares the number of locomotives (bins) used by DP and FFD across various problem instances, illustrating that FFD typically produces solutions close to optimal. Figure 3 presents the runtime comparison, highlighting the efficiency of FFD compared to the significantly higher computation time required by DP as instance size grows. Finally, Figure 4 shows a histogram of approximation ratios (FFD solution over DP optimal) for instances with 14 trains, confirming that FFD rarely exceeds the optimal number of locomotives by more than one.

These results validate FFD as a practical heuristic for large-scale locomotive assignment problems, while DP remains a valuable tool for benchmarking and small instances.




Figure 2: DP Bin vs FFD Bin




Figure 3: Runtime Comparison: DP vs FFD




Figure 4: Histogram of Approximation Ratios (n=14)

6 Python Implementation

Below is a condensed version of the core algorithms. Full versions may be included in an appendix.

```
def ffd(items, B):
    items = sorted(items, reverse=True)
    bins = []
    for x in items:
        placed = False
        for i in range(len(bins)):
            if bins[i] >= x:
                bins[i] -= x
                placed = True
                break
        if not placed:
            bins.append(B - x)
    return len(bins)
```

```

def exact_binpacking(a, B):
    n = len(a)
    N = 1 << n
    INF = 10**9
    dp = [INF]*N
    load = [0]*N
    dp[0] = 1
    for mask in range(N):
        if dp[mask] == INF:
            continue
        for i in range(n):
            if not (mask & (1<<i)):
                nm = mask | (1<<i)
                if load[mask] + a[i] <= B:
                    if dp[nm] > dp[mask]:
                        dp[nm] = dp[mask]
                        load[nm] = load[mask]
                else:
                    if dp[nm] > dp[mask] + 1:
                        dp[nm] = dp[mask] + 1
                        load[nm] = a[i]
    return dp[-1]

```

7 Conclusion

We introduced the Maintenance-Constrained Locomotive Assignment (MCLA) problem, proved strong NP-hardness via 3-Partition reduction, and presented exact and approximate algorithms. Because the problem is structurally equivalent to Bin Packing, classical heuristics like FFD are effective, while ILP solves moderate-scale instances optimally.

This connection between a real-world railway scheduling challenge and a canonical NP-hard problem provides both theoretical insight and practical algorithms for railway operators.

References

- [1] Garey and Johnson. *Computers and Intractability: A Guide to the Theory of NP-Completeness*. Freeman, 1979.

A Python Code for Experiments

```
1 import os
2 import random
3 import time
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7
8 def ffd(items, B):
9     items = sorted(items, reverse=True)
10    bins = []
11    for x in items:
12        placed = False
13        for i in range(len(bins)):
14            if bins[i] >= x:
15                bins[i] -= x
16                placed = True
```

```

17         break
18     if not placed:
19         bins.append(B - x)
20     return len(bins)
21
22 def exact_binpacking(a, B):
23     n = len(a)
24     N = 1 << n
25     INF = 10**9
26     dp = [INF]*N
27     load = [0]*N
28     dp[0] = 1
29     for mask in range(N):
30         if dp[mask] == INF:
31             continue
32         for i in range(n):
33             if not (mask & (1<<i)):
34                 nm = mask | (1<<i)
35                 if load[mask] + a[i] <= B:
36                     if dp[nm] > dp[mask]:
37                         dp[nm] = dp[mask]
38                         load[nm] = load[
39                             mask] + a[i]
40                 else:
41                     if dp[nm] > dp[mask] +
42                         1:
43                         dp[nm] = dp[mask] +
44                             1
45                         load[nm] = a[i]
46
47     return dp[-1]
48
49 def generate_instance(n, B, seed=None):
50     random.seed(seed)
51     # Generate train usages uniformly
52     # between 1 and B/2 to avoid trivial
53     # bins
54     return [random.randint(1, B//2) for _ -
55             in range(n)]
56
57 def run_experiments(out_dir="mcla_experiments", seed=42):
58     os.makedirs(out_dir, exist_ok=True)
59     random.seed(seed)
60     np.random.seed(seed)
61
62     n_values = list(range(8, 18)) # n from
63     # 8 to 17 for exact DP feasibility
64     B = 100
65
66     results = []
67     for n in n_values:
68         for trial in range(10): # 10
69             trials per n
70             items = generate_instance(n, B)
71             # Run exact DP (time it)
72             start = time.time()
73             exact_bins = exact_binpacking(
74                 items, B)
75             exact_time = time.time() -
76                 start
77
78             # Run FFD heuristic (time it)
79
80

```

```

    start = time.time()
    ffd_bins = ffd(items, B)
    ffd_time = time.time() - start

    approx_ratio = ffd_bins /
                    exact_bins if exact_bins >
    0 else np.nan

    results.append({
        "n": n,
        "trial": trial,
        "exact_bins": exact_bins,
        "ffd_bins": ffd_bins,
        "approx_ratio":
            approx_ratio,
        "exact_time": exact_time,
        "ffd_time": ffd_time
    })
}

df = pd.DataFrame(results)
csv_path = os.path.join(out_dir, "experiment_results.csv")
df.to_csv(csv_path, index=False)

# Plot 1: Approximation ratio mean +
# median vs n
plt.figure(figsize=(8,6))
grouped = df.groupby("n")["approx_ratio"]
grouped = grouped.agg(["mean", "median"])
grouped = grouped.reset_index()
grouped = grouped[["n", "mean", "median"]]
grouped = grouped.sort_values("n")
grouped = grouped.reset_index()

x = []
y_mean = []
y_median = []
for n, series in grouped:
    x.append(n)
    y_mean.append(series.mean())
    y_median.append(series.median())
plt.plot(x, y_mean, marker='o', label="Mean ratio")
plt.plot(x, y_median, marker='x', label="Median ratio")
plt.xlabel("Number of trains (n)")
plt.ylabel("Approximation ratio (FFD / Exact)")
plt.title("FFD Approximation Ratio vs Exact DP")
plt.legend()
plt.grid(True)
ratio_vs_n = os.path.join(out_dir, "ratio_vs_n.png")
plt.savefig(ratio_vs_n)
plt.close()

# Plot 2: Approximation ratio boxplot
# per n with fixed labels
plt.figure(figsize=(8,6))
n_values = sorted(df["n"].unique())
data_to_plot = [df[df["n"] == n][
    "approx_ratio"].dropna() for n in
    n_values]
plt.boxplot(data_to_plot, labels=
    n_values, patch_artist=True)
plt.xlabel("Number of trains (n)")

```

```

115     plt.ylabel("Approximation ratio (FFD / Exact)")
116     plt.title("FFD Approximation Ratio Boxplot")
117     plt.grid(True)
118     plt.xticks(ticks=range(1, len(n_values) + 1), labels=n_values)
119     plt.tight_layout()
120     approx_ratio_boxplot = os.path.join(
121         out_dir, "approx_ratio_boxplot.png"
122     )
123     plt.savefig(approx_ratio_boxplot)
124     plt.close()
125
126 # Plot 3: Runtime comparison (mean) vs n
127 plt.figure(figsize=(8,6))
128 grouped_time = df.groupby("n")[["exact_time", "ffd_time"]].mean()
129 plt.plot(grouped_time.index,
130          grouped_time["exact_time"], marker='o', label="Exact DP")
131 plt.plot(grouped_time.index,
132          grouped_time["ffd_time"], marker='x', label="FFD heuristic")
133 plt.xlabel("Number of trains (n)")
134 plt.ylabel("Runtime (seconds)")
135 plt.title("Runtime Comparison")
136 plt.legend()
137 plt.grid(True)
138 runtime_comparison = os.path.join(
139     out_dir, "runtime_comparison.png")
140 plt.savefig(runtime_comparison)
141 plt.close()
142
143 # Plot 4: Histogram of approximation ratios for max n
144 max_n = max(n_values)
145 plt.figure(figsize=(8,6))
146 approx_ratios = df[df["n"] == max_n][["approx_ratio"]].dropna()
147 plt.hist(approx_ratios, bins=10,
148          edgecolor='black')
149 plt.xlabel("Approximation ratio (FFD / Exact)")
150 plt.ylabel("Frequency")
151 plt.title(f"Approximation Ratio Histogram (n={max_n})")
152 plt.grid(True)
153 ratio_histogram = os.path.join(out_dir, "ratio_histogram.png")
154 plt.savefig(ratio_histogram)
155 plt.close()
156
157 # Plot 5: Density sweep (vary B with fixed n)
158 n_fixed = 14
159 B_values = list(range(50, 201, 10))
160 density_results = []
161 for B_curr in B_values:
162     ratios = []
163     for _ in range(5): # 5 trials each

```

```

158     items = generate_instance(
159         n_fixed, B_curr)
160     exact = exact_binpacking(items,
161                             B_curr)
162     ffd_res = ffd(items, B_curr)
163     if exact > 0:
164         ratios.append(ffd_res /
165                         exact)
166     density_results.append(np.mean(
167         ratios) if ratios else np.nan)
168 plt.figure(figsize=(8,6))
169 plt.plot(B_values, density_results,
170          marker='o')
171 plt.xlabel("Capacity B")
172 plt.ylabel("Mean Approximation Ratio")
173 plt.title(f"Density Sweep (n={n_fixed})")
174 plt.grid(True)
175 density_sweep = os.path.join(out_dir, "density_sweep.png")
176 plt.savefig(density_sweep)
177 plt.close()
178
179 print(f"Wrote outputs to: {out_dir}")
180 print(f"CSV: {csv_path}")
181 print("Plots:", ratio_vs_n,
182       approx_ratio_boxplot,
183       runtime_comparison, ratio_histogram
184       , density_sweep)
185
186 if __name__ == "__main__":
187     run_experiments()

```

Listing 1: Experiment scripts for MCLA problem: exact DP, FFD heuristic, data generation, and plotting.

B LLM-Assisted Methodology and Documentation

B.1 Purpose of LLM Usage

Large Language Model (LLM) assistance was used selectively for:

- (1) verifying mathematical soundness of recursive and inductive definitions related to the NP-completeness proof,
- (2) refining \LaTeX syntax for theorems, proofs, and algorithm listings,
- (3) drafting and formatting documentation.

B.2 Representative Professional Prompts and Results

Prompt 1 – Mathematical Structure and Proof Validation:

"Review the NP-completeness reduction from 3-Partition to Maintenance-Constrained Locomotive Assignment (MCLA). Verify the correctness of the inductive proof and formalize it using rigorous \LaTeX notation. Identify and correct any logical flaws in the if-and-only-if arguments."

Result 1:

The model identified subtle gaps in the initial proof, such as clarifying that each locomotive must be fully utilized due to capacity constraints and item size bounds. It proposed a corrected, formal proof of equivalence between 3-Partition and MCLA feasibility,

using induction on the number of locomotives and item sets. The finalized proof appears verbatim in Section 3.

Prompt 2 – Formatting and Documentation:

"Polish all theorem, definition, and proof environments for professional academic presentation. Ensure mathematical symbols use AMS conventions, spacing is consistent, and equations compile without errors. Add meaningful captions to figures and align section structure with IEEE/ACM formatting guidelines."

Result 2:

The LLM returned clean, well-structured \LaTeX fragments utilizing `amsthm` environments, corrected spacing and math mode syntax, and provided IEEE-compliant captions and figure placements. This improved the overall readability and professional appearance of the manuscript.